

Arrays & Strings

Readings: CP:AMA 8.1, 8.3, 9.3, 10, 12.1, 12.2, 12.3, 13

Arrays

The only two types of “compound” data storage *built-in* to C are *structures* and *arrays*.

```
int my_array[6] = {4, 8, 15, 16, 23, 42};
```

An array is a data structure that contains a **fixed number** of elements that all have the **same type**.

Because arrays are *built-in* to C, they are used for many tasks where *lists* are used in Racket, but **arrays and lists are very different**. In Section 11 we construct Racket-like lists in C.

```
int my_array[6] = {4, 8, 15, 16, 23, 42};
```

To define an array we must know the **length** of the array **in advance** (we address this limitation in Section 10).

Each individual value in the array is known as an *element*. To access an element, its *index* is required.

The first element of `my_array` is at index 0, and it is written as `my_array[0]`. The second element is `my_array[1]`. The last element is `my_array[5]`.

In computer science we often start counting at 0.

example: accessing array elements

Each individual array element can be used as if it was a variable.

```
int a[6] = {4, 8, 15, 16, 23, 42};

int j = a[0];           // j is 4
int k = a[j];           // k is 23
int *p = &a[k-20];      // p points at a[3]

a[2] = a[a[0]];         // a[2] is now 23
a[0]++;                 // a[0] is now 5
++a[1];                 // a[1] is now 9
```

example: arrays & iteration

Arrays and iteration are a powerful combination.

```
int a[6] = {4, 8, 15, 16, 23, 42};  
int sum = 0;
```

```
for (int i = 0; i < 6; ++i) {  
    printf("a[%d] = %d\n", i, a[i]);  
    sum += a[i];  
}  
printf("sum = %d\n", sum);
```

```
a[0] = 4  
a[1] = 8  
a[2] = 15  
a[3] = 16  
a[4] = 23  
a[5] = 42  
sum = 108
```

Array initialization

Like variables, an uninitialized array

```
int a[5];
```

is zero-filled if the array is *global*. If the array is *local*, it is filled with arbitrary (“garbage”) values from the stack.

As with structures, the array braces (`{}`) syntax is only valid in **initialization**.

If there are not enough elements in the braces, the remaining values are initialized to zero (even with local arrays).

```
int b[5] = {1, 2, 3};    // b[3] & b[4] = 0
int c[5] = {0};          // c[0]...c[4] = 0
```

If an array is initialized, the length of the array can be omitted from the declaration and *automatically* determined from the number of elements in the initialization.

```
int a[] = {4, 8, 15, 16, 23, 42}; // int a[6] = ...
```

This syntax is only allowed if the array is initialized.

```
int b[]; // INVALID
```

Similar to structures, C99 supports a partial initialization syntax.

```
int a[100] = { [50] = 1, [25] = -1, [75] = 3 };
```

Omitted elements are initialized to zero.

C99 allows the length of an **uninitialized local array** to be determined *while the program is running*. The size of the stack frame is increased accordingly.

```
int count;  
printf("How many numbers? ");  
scanf("%d",&count);
```

```
int a[count];           // count determined at run-time
```

This approach has many disadvantages and in the most recent version of C (C11), this feature was made optional. In Section 10 we see a better approach.

Array size

The **length** of an array is the number of elements in the array.

The **size** of an array is the number of bytes it occupies in memory.

An array of k elements, each of size s , requires exactly $k \times s$ bytes.

In the C memory model, array elements are adjacent to each other.

Each element of an array is placed in memory immediately after the previous element.

If `a` has six elements (`int a[6];`) the size of `a` is
 $(6 \times \text{sizeof}(\text{int})) = 6 \times 4 = 24$.

Not everyone uses the same terminology for length and size.

example: array in memory

```
int a[6] = {4, 8, 15, 16, 23, 42};  
printf("&a[0] = %p ... &a[5] = %p\n", &a[0], &a[5]);  
&a[0] = 0x5000 ... &a[5] = 0x5014
```

addresses	contents (4 bytes)
0x5000 ... 0x5003	4
0x5004 ... 0x5007	8
0x5008 ... 0x500B	15
0x500C ... 0x500F	16
0x5010 ... 0x5013	23
0x5014 ... 0x5017	42

Array length

C does not explicitly keep track of the array **length** as part of the array data structure.

You must keep track of the array length separately.

Often, the array length is stored in a separate variable.

```
const int a_length = 6;  
int a[a_length] = {4, 8, 15, 16, 23, 42};
```

```
const int a_length = 6;  
int a[a_length];
```

The above definition is fine in `Seashell`, but some C environments do not allow the length of the array to be specified by a variable.

In those environments, the `#define` syntax is more often used. This is common in CP:AMA.

```
#define A_LENGTH 6  
int a[A_LENGTH];
```

Theoretically, in some circumstances you could use `sizeof` to determine the length of an array.

```
int len = sizeof(a) / sizeof(a[0]);
```

The CP:AMA textbook uses this on occasion.

However, in practice, this should never be done, as the `sizeof` operator only properly reports the array size in some circumstances.

The array identifier

The value of an array identifier (the array name) can be used as a pointer to the first element of the array.

```
int a[] = {4, 8, 15, 16, 23, 42};  
printf("%p %p %p \n", a, &a, &a[0]);  
printf("%d %d\n", a[0], *a);
```

0x5000 0x5000 0x5000

4 4

The **value** of `a` is the same as the **address** of the array (`&a`), which is also the address of the first element (`&a[0]`).

Dereferencing the array (`*a`) is equivalent to referencing the first element (`a[0]`).

example: array identifier

In an expression, `*a` is the same as `a[0]`.

```
int a[3] = {0, 0, 0};  
printf("a[0] = %d\n", a[0]);  
*a = 5;  
printf("a[0] = %d\n", a[0]);  
  
a[0] = 0;  
a[0] = 5;
```

The array identifier itself cannot be changed or assigned to and is effectively “constant”.

```
int a[3] = {0, 0, 0};  
int b[3] = {1, 2, 3};  
a = b;                                // INVALID
```

Passing arrays to functions

When an array is passed to a function only the **address** of the array is copied into the stack frame. This is more efficient than copying the entire array to the stack.

Typically, the length of the array is unknown, and is provided as a separate parameter.

example: array parameters

```
int sum_array(int a[], int len) {  
    int sum = 0;  
    for (int i = 0; i < len; ++i) {  
        sum += a[i];  
    }  
    return sum;  
}
```

```
int main(void) {  
    int my_array[6] = {4, 8, 15, 16, 23, 42};  
    int sum = sum_array(my_array, 6);  
}
```

Note the parameter syntax: `int a[]`

and the calling syntax: `sum_array(my_array, 6)`.

As we have seen before, passing an address to a function allows the function to change (mutate) the contents at that address.

```
void array_add1(int a[], int len) {  
    for (int i = 0; i < len; ++i) {  
        ++a[i];  
    }  
}
```

It's good style to use the `const` keyword to prevent mutation and communicate that no mutation occurs.

```
int sum_array(const int a[], int len) {  
    int sum = 0;  
    for (int i = 0; i < len; ++i) {  
        sum += a[i];  
    }  
    return sum;  
}
```

Because a structure can contain an array:

```
struct mystruct {  
    int big[1000];  
};
```

It is *especially* important to pass a pointer to such a structure, otherwise, the **entire array** is copied to the stack frame.

```
int slower(struct mystruct s) {  
    ...  
}  
  
int faster(struct mystruct *s) {  
    ...  
}
```

Pointer arithmetic

We have not yet discussed any *pointer arithmetic*.

C allows an integer to be added to a pointer, but the result may not be what you expect.

If p is a pointer, the value of $(p+1)$ **depends on the type** of the pointer p .

$(p+1)$ adds the `sizeof` whatever p points at.

According to the official C standard, pointer arithmetic is only valid **within an array** (or a structure) context. This becomes clearer later.

Pointer arithmetic rules

- When adding an integer `i` to a pointer `p`, the address computed by `(p + i)` in C is given in “normal” arithmetic by:

$$p + i \times \text{sizeof}(*p).$$

- Subtracting an integer from a pointer `(p - i)` works in the same way.
- Mutable pointers can be incremented (or decremented).
`++p` is equivalent to `p = p + 1`.

- You cannot add two pointers.
- You can subtract a pointer `q` from another pointer `p` if the pointers are the same type (point to the same type). The value of `(p - q)` in C is given in “normal” arithmetic by:

$$(p - q) / \text{sizeof}(*p).$$

In other words, if `p = q + i` then `i = p - q`.

- Pointers (of the same type) can be compared with the comparison operators: `<`, `<=`, `==`, `!=`, `>=`, `>` (e.g., `if (p < q) . . .`).

Pointer arithmetic and arrays

Pointer arithmetic is useful when working with **arrays**.

Recall that for an array `a`, the value of `a` is the address of the first element (`&a[0]`).

Using pointer arithmetic, the address of the second element `&a[1]` is `(a + 1)`, and it can be referenced as `*(a + 1)`.

The array indexing syntax (`[]`) is an **operator** that performs *pointer arithmetic*.

`a[i]` is *equivalent* to `*(a + i)`.

In ***array pointer notation***, square brackets (`[]`) are not used, and all array elements are accessed through pointer arithmetic.

```
int sum_array(const int *a, int len) {  
    int sum = 0;  
    for (const int *p = a; p < a + len; ++p) {  
        sum += *p;  
    }  
    return sum;  
}
```

Note that the above code behaves **identically** to the previously defined `sum_array`:

```
int sum_array(const int a[], int len) {  
    int sum = 0;  
    for (int i = 0; i < len; ++i) {  
        sum += a[i];  
    }  
    return sum;  
}
```


another example: pointer notation

```
// count_match(item, a, len) counts the number of  
// occurrences of item in the array a
```

```
int count_match(int item, const int *a, int len) {  
    int count = 0;  
    const int *p = a;  
    while (p < a + len) {  
        if (*p == item) {  
            ++count;  
        }  
        ++p;  
    }  
    return count;  
}
```

The choice of notation (pointers or []) is a matter of style and context. You are expected to be comfortable with both.

C makes no distinction between the following two function headers (declarations):

```
int sum_array(const int a[], int len) {...}    // a[]  
int sum_array(const int *a, int len) {...}    // *a
```

In most contexts, there is no practical difference between an array identifier and a (*constant*) pointer.

The subtle differences between an array and a pointer are discussed at the end of this Section.

Multi-dimensional data

All of the arrays seen so far have been one-dimensional (1D) arrays.

We can represent multi-dimensional data by “mapping” the higher dimensions down to one.

For example, consider a 2D array with 2 rows and 3 columns.

```
1 2 3
7 8 9
```

We can represent the data in a simple one-dimensional array.

```
int data[6] = {1, 2, 3, 7, 8, 9};
```

To access the entry in row `r` and column `c`, we simply access the element at `data[r*3 + c]`.

In general, it would be `data[row * NUMCOLS + col]`.

C supports multiple-dimension arrays, but they are not covered in this course.

```
int two_d_array[2][3];  
int three_d_array[10][10][10];
```

When multi-dimensional arrays passed as parameters, the second (and higher) dimensions must be fixed.

(e.g., `int function_2d(int a[][10], int numRows)`).

Internally, C represents a multi-dimensional array as a 1D array and performs “mapping” similar to the method described in the previous slide.

See CP:AMA sections 8.2 & 12.4 for more details.

Abstract array functions

In Racket, **Abstract List Functions (ALFs)** (e.g., `filter`, `map`, `foldl`) are powerful and flexible tools for processing lists.

In C, we can create ***Abstract Array Functions (AAFs)***.

```
int add1(int n) {  
    return n + 1;  
}  
  
int main(void) {  
    int a[] = {4, 8, 15, 16, 23, 42};  
    print_array(a, 6);  
    array_map(add1, a, 6);    // we need to define this  
    print_array(a, 6);  
}
```

4 8 15 16 23 42
5 9 16 17 24 43

Function pointers

In Racket, functions are *first-class values*. For example, Racket functions are values that can be stored in variables and data structures, passed as arguments and returned by functions.

In C, functions are not first-class values. However, all of the aforementioned things can be done with ***function pointers***.

A *function pointer* stores the starting address of a function within the code section. The value of a function identifier is its starting address.

A significant difference is that new Racket functions can be created during program execution, while in C they cannot. A function pointer can only point to a function that already exists.

The declaration for a function pointer includes the *return type* and all of the *parameter types*, which makes them a little messy.

For example, consider the following function pointer `fp` that points at the function `add1`:

```
int add1(int i) {  
    return i + 1;  
}
```

```
int (*fp)(int) = add1;
```

The syntax to declare a function pointer with name `fp` is:

```
return_type (*fp)(param1_type, param2_type, ...)
```

examples: function pointer declarations

```
int functionA(int i) {...}  
int (*fpA)(int) = functionA;
```

```
char functionB(int i, int j) {...}  
char (*fpB)(int, int) = functionB;
```

```
int functionC(int *ptr, int i) {...}  
int (*fpC)(int *, int) = functionC;
```

```
int *functionD(int *ptr, int i) {...}  
int *(*fpD)(int *, int) = functionD;
```

```
struct posn functionE(struct posn *p, int i) {...}  
struct posn (*fpE)(struct posn *, int) = functionE;
```

In an exam, we would not expect you to remember the syntax for declaring a function pointer.

Array map

Aside from the function pointer parameter syntax, the definition of `array_map` is straightforward.

```
// effects: replaces each element a[i] with f(a[i])
```

```
void array_map(int (*f)(int), int a[], int len) {  
    for (int i=0; i < len; ++i) {  
        a[i] = f(a[i]);  
    }  
}
```

example: Array map

```
#include "array_map.h"
```

```
int add1(int i) { return i + 1; }  
int sqr(int i) { return i * i; }  
int print_element(int i) {  
    printf("%d ", i); return i;  
}
```

```
int main(void) {  
    int a[] = {4, 8, 15, 16, 23, 42};  
    array_map(print_element, a, 6); printf("\n");  
    array_map(add1, a, 6);  
    array_map(print_element, a, 6); printf("\n");  
    array_map(sqr, a, 6);  
    array_map(print_element, a, 6); printf("\n");  
}
```

4 8 15 16 23 42

5 9 16 17 24 43

25 81 256 289 576 1849

Alternatively, the function passed to `array_map` could be a `void` function that accepts a pointer to each element.

```
void add1(int *p) {
    *p += 1;
}

void alt_array_map(void (*f)(int *), int a[], int len) {
    for (int i=0; i < len; ++i) {
        f(a + i);                // same as f(&a[i])
    }
}
```

Our `array_map` only works with arrays of integers.

In practice, a generic AAF would likely be implemented with `void` pointers, which are described in Section 12.

Array foldl

```
// foldl(f, base, a, len) returns:  
//  f(a[len-1], f(a[len-2], ..., f(a[1], f(a[0], base))))  
  
int array_foldl(int (*f)(int, int), int base, int a[],int len) {  
    for (int i=0; i < len; ++i) {  
        base = f(a[i], base);  
    }  
    return base;  
}  
  
// the only change for foldr would be a different for loop:  
//  for (int i=len-1; i >= 0; --i) {
```

In Section 10 we introduce dynamic arrays, where a `filter` AAF makes more sense.

Because C does not have anonymous functions or the ability to dynamically generate closure functions (e.g., `lambda`), AAFs are not as convenient or as popular as ALFs in Racket.

In practice, it is not much more effort to write a function to perform an action on an entire array instead of a single element.

```
int add1(int i) {  
    return i + 1;  
}  
  
void array_add1(int a[], int len) {  
    for (int i=0; i < len; ++i) {  
        a[i]++;  
    }  
}
```

Strings

There is no built-in C *string* type. The “**convention**” is that a C string is an **array of characters**, terminated by a *null character*.

```
char my_string[4] = {'c', 'a', 't', '\\0'};
```

The *null character*, also known as a null *terminator*, is a `char` with a value of zero. It is often written as `'\\0'` instead of just `0` to improve communication and indicate that a null character is intended.

`'\\0'` (ASCII 0) is different than `'0'` (ASCII 48), which is the character for the symbol zero.

String initialization

In addition to the regular array initialization syntax, `char` arrays also support a double quote (") notation. When combined with the automatic size declaration ([]), the size includes the null terminator.

The following definitions create equivalent 4-character arrays:

```
char a[] = {'c', 'a', 't', '\0'};  
char b[] = {'c', 'a', 't', 0};  
char c[4] = {'c', 'a', 't'};  
char d[] = { 99, 97, 116, 0};  
char e[4] = "cat";  
char f[] = "cat";
```

This array **initialization** notation is **different** than the double quote notation used in expressions (e.g., in `printf("string")`).

String literals

The C strings used in statements (e.g., with `printf` and `scanf`) are known as *string literals*.

```
printf("i = %d\n", i);  
printf("the value of j is %d\n", j);
```

For each string literal, a null-terminated `const char` array is created in the *read-only data* section.

In the code, the occurrence of the *string literal* is replaced with address of the corresponding array.

The “*read-only*” section is also known as the “*literal pool*”.

example: string literals

```
void foo(int i, int j) {  
    printf("i = %d\n", i);  
    printf("the value of j is %d\n", j);  
}
```

Although no array name is actually given to each literal, it is helpful to imagine that one is:

```
const char foo_string_literal_1[] = "i = %d\n";  
const char foo_string_literal_2[] = "the value of j is %d\n";  
  
void foo(int i, int j) {  
    printf(foo_string_literal_1, i);  
    printf(foo_string_literal_2, j);  
}
```

You should not try to modify a string literal. The behaviour is undefined, and it causes an error in Seashell.

Null termination

Because strings are null terminated, we do not have to pass the array length to every function.

```
// e_count(s) counts the # of e's in string s
```

```
int e_count(const char *s) {  
    int count = 0;  
    while (*s) { // not the null terminator  
        if ((*s == 'e') || (*s == 'E')) ++count;  
        ++s;  
    }  
    return count;  
}
```

As with “regular” arrays, it is good style to have `const` parameters to communicate that no changes (mutation) occurs to the string.

strlen

The `string` library (`#include <string.h>`) provides many useful functions for processing strings (more on this library later).

The `strlen` function returns the length of the *string*, **not** necessarily the length of the *array*. It does **not include** the null character.

```
int my_strlen(const char *s) {  
    int len = 0;  
    while (s[len]) {  
        ++len;  
    }  
    return len;  
}
```

Here is an alternative implementation of `my_strlen` that uses pointer arithmetic.

```
int my_strlen(const char *s) {  
    const char *p = s;  
    while (*p) {  
        ++p;  
    }  
    return (p - s);  
}
```

Lexicographical order

Characters can be easily compared ($c1 < c2$) as they are numbers, so the character **order** is determined by the ASCII table.

If we try to compare two strings ($s1 < s2$), C compares their *pointers*, which is not helpful.

To compare strings we are typically interested in using a **lexicographical order**.

Strings require us to be more careful with our terminology, as “smaller than” and “greater than” are ambiguous: are we considering just the **length** of the string? To avoid this problem we use **precedes** (“before”) and **follows** (“after”).

To compare two strings using a **lexicographical order**, we first compare the first character of each string. If they are different, the string with the smaller first character *precedes* the other string. Otherwise (the first characters are the same), the second characters are compared, and so on.

If the end of one string is encountered, it *precedes* the other string. Two strings are equal (the same) if they are the same length and all of their characters are identical.

The following strings are in lexicographical order:

" " "a" "az" "c" "cab" "cabin" "cat" "catastrophe"

The `<string.h>` library function `strcmp` uses lexicographical ordering.

`strcmp(s1, s2)` returns zero if the strings are identical. If `s1` precedes `s2`, it returns a negative integer. Otherwise (`s1` follows `s2`) it returns a positive integer.

```
int my_strcmp(const char *s1, const char *s2) {
    while (*s1 == *s2) {
        if ((*s1 == '\0') && (*s2 == '\0')) return 0;
        ++s1;
        ++s2;
    }
    if (*s1 < *s2) return -1;
    return 1;
}
```

To compare if two strings are *equal* (identical), use the `strcmp` function.

The equality operator (`==`) only compares the *addresses* of the strings, and not the contents of the arrays.

```
char a[] = "the same?";  
char b[] = "the same?";  
char *s = a;  
  
if (a == b) ...           // False (diff. addresses)  
if (strcmp(a,b) == 0) ... // True  (proper comparison)  
if (a == s) ...           // True  (same addresses)
```


Lexicographical orders can be used to compare (and sort) any *sequence* of elements (arrays, lists, ...) and not just strings.

The following Racket function lexicographically compares two lists of numbers:

```
(define (lon<=? lon1 lon2)
  (cond [(empty? lon1) #t]
        [(empty? lon2) #f]
        [(< (first lon1) (first lon2)) #t]
        [(< (first lon2) (first lon1)) #f]
        [else (lon<=? (rest lon1) (rest lon2))]))
```

```
(lon<=? '(4 9 1 2 1) '(4 5 9)) ; => #f
(lon<=? '(4 3) '(4 3 2))       ; => #t
```

String I/O

The `printf` placeholder for strings is `%s`.

```
char a[] = "cat";  
printf("the %s in the hat\n", a);
```

`printf` prints out characters until the null character is encountered.

When using `%s` with `scanf`, it stops reading the string when a “white space” character is encountered (e.g., a space or `\n`).

`scanf("%s")` is useful for reading in one “word” at a time.

```
char name[81];  
printf("What is your first name? ");  
scanf("%s", name);
```

You must be very careful to reserve enough space for the string to be read in, and **do not forget the null character**.

In this example, the array is 81 characters and can accommodate first names with a length of up to 80 characters.

What if someone has a *really* long first name?

example: scanf

```
int main(void) {
    char command[8];
    int balance = 0;
    while (1) {
        printf("Command? ('balance', 'deposit', or 'q' to quit): ");
        scanf("%s", command);
        if (strcmp(command, "balance") == 0) {
            printf("Your balance is: %d\n", balance);
        } else if (strcmp(command, "deposit") == 0) {
            printf("Enter your deposit amount: ");
            int dep;
            scanf("%d", &dep);
            balance += dep;
        } else if (strcmp(command, "q") == 0) {
            printf("Bye!\n"); break;
        } else {
            printf("Invalid command. Please try again.\n");
        }
    }
}
```

In this banking example, entering a long command causes C to write characters beyond the size of the `command` array. Eventually, it overwrites the memory where `balance` is stored.

This is known as a ***buffer overrun*** (or *buffer overflow*). The C language is especially susceptible to *buffer overruns*, which can cause serious stability and security problems.

In this introductory course, having an appropriately sized array and using `scanf` is “good enough”.

In practice you would **never** use this insecure method for reading in a string.

The `gets` function does not stop when a space is encountered, and reads in characters until a newline (`\n`) is encountered. It is also very susceptible to overruns, but is convenient to use in this course.

```
char name[81];  
printf("What is your full name? ");  
gets(name);
```

There are C library functions that are more secure than `scanf` and `gets`.

A popular strategy to avoid overruns is to only read in one character at a time (e.g., with `scanf("%c")` or `getchar`). For an example of using `getchar` to avoid overruns, see CP:AMA 13.3.

Two additional `<string.h>` library functions that are useful, but susceptible to buffer overruns are:

`strcpy(char *dest, const char *src)` overwrites the contents of `dest` with the contents of `src`.

`strcat(char *dest, const char *src)` copies (appends or concatenates) `src` to the end of `dest`.

You should always ensure that the `dest` array is large enough (and don't forget the null terminator).

With this simple implementation of `strcpy`:

```
char *strcpy(char *dst, const char *src) {  
    char *d = dst;  
    do {  
        *d = *src;  
        ++d; ++s;  
    } while (*src);  
    return dst;  
}
```

you can crash your program:

```
char c[] = "spam";  
strcpy(c + 4, c);
```

Because the start of the destination is also the null terminator of the source, the source never terminates and it fills up the memory with `spamspamspace` . . . until a crash occurs.

While *writing* to a buffer can cause dangerous buffer overruns, *reading* an improperly terminated string can also cause problems.

```
char c[3] = "cat";    // NOT properly terminated!
printf("%s\n", c);
printf("The length of c is: %d\n", strlen(c));

cat????????????????
The length of c is: ??
```

The string library has “safer” versions of many of the functions that stop when a maximum number of characters is reached.

For example, `strnlen`, `strncmp`, `strncpy` and `strncat`.

Arrays vs. pointers

Earlier, we said arrays and pointers are *similar* but **different**.

Consider the following two string definitions:

```
void f(void) {  
    char a[] = "pointers are not arrays";  
    char *p  = "pointers are not arrays";  
    ...  
}
```

- The first reserves space for an initialized 24 character array (**a**) in the stack frame (24 bytes).
- The second reserves space for a **char** pointer (**p**) in the stack frame (8 bytes), *initialized* to point at a string literal (**const char** array) created in the read-only data section.

example: more arrays vs. pointers

```
char a[] = "pointers are not arrays";  
char *p = "pointers are not arrays";  
char d[] = "different string";
```

`a` is a `char` array. The *identifier* `a` has a constant value (the address of the array), but the elements of `a` can be changed.

```
a = d;           // INVALID  
a[0] = 'P';      // VALID
```

`p` is a `char` pointer. `p` is initialized to point at a string literal, but `p` can be changed to point at any `char`.

```
p[0] = 'P';      // INVALID (p points at a const literal)  
p = d;           // VALID  
p[0] = 'D';      // NOW VALID (p points at d)
```

An array is very similar to a **constant** pointer.

```
int a[6] = {4, 8, 15, 16, 23, 42};  
int * const p = a;
```

In most practical expressions `a` and `p` would be equivalent. The only significant differences between them are:

- `a` is the same as `&a`, while `p` and `&p` have different values
- `sizeof(a)` is 24, while `sizeof(p)` is 8

`p` also requires an additional four bytes of storage.

Goals of this Section

At the end of this section, you should be able to:

- define and initialize arrays and strings
- use iteration to loop through arrays
- use pointer arithmetic
- explain how arrays are represented in the memory model, and how the array index operator (`[]`) uses pointer arithmetic to access array elements in constant time

- use both array index notation ([]) and array pointer notation and convert between the two
- represent multi-dimensional data in a single-dimensional array
- explain and demonstrate the use of the null termination convention for strings
- explain string literals and the difference between defining a string array and a string pointer
- sort a string or sequence lexicographically

- use I/O with strings and explain the consequences of buffer overruns
- use `<string.h>` library functions (when provided with a well documented interface)